

CSC111 Project Topic: A Tool to Analyze Consequences of Road Segment Closure by Graphical Analysis of Ontario's Road Network

Khalid Alshikhmubarak, Elias Mohammed A Alhazaa, Fayiz Ali

March 30, 2026

Problem Description and Research Question

Transportation networks rely on interconnected highways, roads, bridges, and tunnels. These networks are always vulnerable to construction closures, accidents, and natural disasters, which could affect the structure or even connectivity of the road network. Each year, numerous cities are threatened by natural disasters that could temporarily close down significant routes. When these routes become unusable, people may not have access to fundamental necessities, such as hospitals and food providers. Therefore, it is important to study transportation networks to prevent an area from being isolated from the rest of its city during emergencies, or even if an isolation does not happen, then to understand how deep would the impact of a road segment closure for emergency/construction reasons be on motorists. This predictive analysis would help prepare the province better and take proactive measures to minimize the impacts of road closures. This goal motivates us to develop a tool to deeply analyze road networks: in particular to use technology and computer science to analyze how road closure would impact travel.

The province of Ontario is the most populous province in Canada, and about half of its population lives in the Greater Toronto Area (Ontario Legislative Assembly). Thus, it is imperative that it keeps a check on the impediments people face due to road closures. A week ago, there was a closure on College and Bay Streets in the Downtown of Toronto, which forced drivers and public transport vehicles to take another route. Additionally, each year, Ontario experiences a variety of hazardous weather events that close roads. For instance, just 3 months ago in January 2026, the snowstorm closed down main highways like Don Valley Parkway and a section of Highway 400 (The Weather Network). **The goal of this project is to create a tool that helps transportation people in charge analyze and understand how much a closure of particular swaths of road would disrupt commute and motorists in Ontario by shortest path graphical analysis.**

To be more specific, the analysis by this tool would be based on the extracting and displaying the next shortest paths (not just 1, but as many equivalent shortest paths) to get from the two points that the set of roads that closed down used to connect. As length of roads are usually with decimal precision, therefore, it is hopeless to expect that there would ever be two shortest paths to those two junctions. For instance, consider how a shortest path of 1.2 meters would be considered as not equivalent to a path of 1.3 meters to get from point A to point B, even though they both just differ by 0.1 meters. However, if we ignore the decimals, which we can do because their contribution to length of road is trivial to even pedestrians, just like in previous example where difference of a path not being a shortest path was just 0.1 meters. Hence, our analysis this way would give us a more realistic view of what paths people will choose to go from junction A to junction B, after the road has closed down between those two junctions. Highlighting the length of new shortest path side by side with length of previous shortest path before road closure would give an idea of the scale of impact and that too from different point of views. For instance, if the distance of shortest path goes from 100 meters to 1000 meters, that's a massive problem for pedestrians (but not for motorists). Lastly, actually highlighting the new shortest paths on map one by one would give the person in charge an idea of the structure of roads. This means if a if one, instead of going straight from junction A to junction B, now has to because of road closures take two left turns, then that would compound the time for not only the motorists who used to use that path, but also the motorists who have to now wait longer in queue as there are now other cars waiting for left turn on the signal. Ambulances who have to use that route will also take longer due to more cars using the same route and congesting at intersections. Therefore, showing different shortest paths side by side, previous shortest path length before closure vs after closure and ignoring small trivialities like decimals would give a great idea for the person responsible about the impacts of a future road closure. This would then help that person take proactive decisions to mitigate the effects of road closures.

The dataset to be used is the Ontario Road Network from Ontario Ministry of Natural Resources, available at <https://geohub.lio.gov.on.ca/datasets/mnrf::ontario-road-network-orn-road-net-element/about> (Ontario Ministry of Natural Resource). As of March 2, 2026 11:54 pm, this dataset contains 1,119,553 edges (roads) where a bidirectional road is counted as two different edges. It also contains 433,478 unique nodes. The following piece of JSON, with similar fields, whose example has been demonstrated already in it, and unneeded metadata in response being omitted for brevity, obtained by hitting one of paginated api https://ws.lioservices.lrc.gov.on.ca/arcgis2/rest/services/LIO_OPEN_DATA/LIO_Open09/MapServer/0/query?outFields=OGF_ID,FROM_JUNCTION_ID,TO_JUNCTION_ID,LENGTH,DIRECTION_OF_TRAFFIC_FLOW,OBJECTID&returnGeometry=true&outSR=4326&f=json&where=OBJECTID%20%3C%2047770000 is:

```
{
  ...
  "features": [
    ...
    {
      "attributes": {
        "OGF_ID": 1500457262,
        "FROM_JUNCTION_ID": 7221681,
        "TO_JUNCTION_ID": 7221677,
        "LENGTH": 434.953,
        "DIRECTION_OF_TRAFFIC_FLOW": "Negative",
        "OBJECTID": 47769490
      },
      "geometry": ...
    },
    {
      "attributes": {
        "OGF_ID": 1500452584,
        "FROM_JUNCTION_ID": 1500124587,
        "TO_JUNCTION_ID": 1500124588,
        "LENGTH": 52.924,
        "DIRECTION_OF_TRAFFIC_FLOW": "Both",
        "OBJECTID": 47769586
      },
      "geometry": {
        "paths": [
          [
            [-78.6984056, 45.0490474990568],
            [-78.6981341, 45.0494830990568]
          ]
        ]
      }
    },
    {
      "attributes": {
        "OGF_ID": 1500452564,
        "FROM_JUNCTION_ID": 6052867,
        "TO_JUNCTION_ID": 6052870,
        "LENGTH": 85.19,
        "DIRECTION_OF_TRAFFIC_FLOW": "Positive",
        "OBJECTID": 47769570
      },
      "geometry": ...
    }
  ]
  ...
}
```

```
]
}
```

Extracted from documentation of this API, the following information about this JSON is important to note (Ontario Ministry of Natural Resources 16, 59). The JSON here shows that road with object id 47769570 runs from vertex id 6052867 (FROM_JUNCTION_ID) to vertex id 6052870 (TO_JUNCTION_ID) only in the forward direction. In contrast, the road from vertex id 1500124587 to vertex id 1500124588 runs both ways, while the road from vertex id 7221681 to 7221677 only runs in reverse direction. The first element of attribute **paths** in **geometry** gives a list of coordinate tuples, that taken together give a real world polyline that connects these two vertices and gives the road its actual shape.

However, the API is paginated and only gives about 2000 records at most for one call. This means 500 API calls would be needed, along with a slight delay as well to prevent the API server from thinking that we are doing a DOS attack on it and hence blocking us. Therefore, we wrote code to hit this API with slight delay and with paginated cursors due to its 2000 roads limit, accumulate all data from it and save it to a file called `ontario_road_network.geojson` (the file shared). Also, we made sure that the url query is in a way that unnecessary data does not get returned: only the fields `OGD_ID`, `FROM_JUNCTION_ID`, `TO_JUNCTION_ID`, `LENGTH`, `GEOMETRY`, `OBJECTID` and `DIRECTION_OF_TRAFFIC_FLOW` were included in our api query and what we saved to the file. To prevent file from becoming huge in size, it does not have indented json. Also, it is in format geojson, which is just a type of json with strict structure imposed on it to make it suitable for geographical computations.

1 Computational Overview

The road network is represented as a weighted directed graph, where each vertex represents an intersection or road endpoint in Ontario, and each edge represents a road segment connecting two of those vertices with a weight representing the distance between them in meters. This graph structure directly models the road network because roads inherently connect two points, and removing an edge in the graph mirrors a real-world road/segment of road closure: allowing the analysis of which road path (next shortest path) that closure would lead the motorists to take to travel between the two vertices this segment used to connect. Graph also allows us to run Dijkstra's algorithm (more on that later), thus allowing computation of shortest paths. Therefore, Graph is the backbone of our project, and it is represented by Graph class in graph.py file

However, Graph alone does not entirely suffice. Because we are computing multiple shortest paths of same length (ignoring decimal part) from a source junction to a target junction, and the usage of trees helps us store and keep track of these paths. The class PathTree, found in path_tree.py, is a tree class to represent multiple paths from the source node to the target node. The target node is the root of the tree, and the leaves are the source nodes, so each path is traced from a leaf (a source node) all the way up to the root (the target node) of the tree. The reason it is upside down is because in our version of Dijkstra's, it is a lot easier and efficient to track each shortest path back from target vertex to source then to do it the other way. PathTree was indispensable for this job, because it is not just about multiple shortest paths to target, but rather a lot of times there are also multiple equivalent shortest paths to another vertex which lies along the shortest path, thus giving the task of tracing all these paths a branching structure. Tree is a natural choice for this representation of multiple paths therefore. However, we don't return a tree back to our other classes, as it is cumbersome to traverse a tree to get a shortest path to display. Rather, after running Dijkstra's algorithm to compute all the shortest paths from a source node to a target node, the returned object is an instance of the ShortestPathResult class in graph.py. This class is used to store all the possible paths from the source to the target node in its instance attribute, `all_shortest_paths` list. This list is populated by calling the function `get_all_possible_paths` method of class PathTree, which returns all possible paths after the PathTree tree has been built by Dijkstra algorithm in the form of a list of list by storing each traversal from depth point of view from root to leaf in a separate list. All paths in this list have the shortest length from the source to the target node, which is stored in the instance attribute of the ShortestPathResult class, `length`. This is done in an optimized manner, as every root in PathTree and its subtrees has a tuple consisting of vertex id and road id that this vertex uses this connect to the parent (useful in case there are multiple direct roads of different length between two neighbouring vertices). A list is not used in our Dijktras algorithm to store a history of shortest paths directly, and rather this PathTree is used, as using a list all the way throughout Dijktras instead of storing just the immediate vertex which is part of shortest path in distance dict (see distance dict in Graph.py for more details) would have slowed it a lot.

Overall, Graph class is the backbone of our project, but it cannot do its job without PathTree.

Types of computations performed on the geojson data:

1. Displaying a graph representation of road network on Ontario map.
2. Making road segment(s) removable (soft delete) by clicking on it/them, using a modified version of Dijkstra's algorithm to compute all shortest paths of equal length (ignoring decimals for reasons discussed earlier) between point A and point B that the removed road segment(s) used to connect previously. Dijkstra's algorithm is used to give one shortest path from one source vertex to all other ones, and it has a time complexity of $O(V + E \log V)$ where E is the number of edges (roads) and V is the number of vertices (intersections/road start/end point), so it is best suited to solve the problem efficiently for this data (GeeksforGeeks). Our modified version gives all shortest paths.
3. Using the same algorithm as above, building a collection of the shortest paths from the graph and highlighting these shortest paths on the map in a distinct color. If no path exists, the program displays a message indicating those two junctions/vertices are now disconnected.
4. Filtering and displaying only the part of graph which is currently in view, as drawing all polylines at once will lead to performance issues. Roads are filtered based on two criteria: whether any coordinate of the road falls within the current map bounds, and whether the road's length exceeds a threshold derived from the current zoom level. At zoom levels below 13, only roads longer than this threshold are shown. At zoom level 13 and above, all roads within bounds are shown. However, all roads present in currently displayed shortest path are displayed in the map regardless of any zoom or bounds
5. Restoring graph state to original one, so previous computations/deletions don't impact the future and fresh ones.

Some notes on architecture of program:

The **Streamlit** library's `st.session_state` is used in `main.py` to persist the **StreamlitManager** instance across reruns. This way, `main.py` acts as a bootstrapping file only, with Streamlit related stuff mostly done in **StreamlitManager** class (Streamlit by default is skewed towards non object oriented code, so this needed bootstrapping this way). The **StreamlitManager** class' display method is what gets called when Streamlit runs the file from top to bottom, while the constructor only runs once, thus providing a convenient and familiar interface in this architecture designed by us to use Streamlit without it getting out of hand and making troubleshooting for bugs extremely tough.

StreamlitManager is our Boundary class, meaning it has absolutely no business logic and has its only job as displaying UI and handling logic related to the display of UI, and handling callbacks.

The **RoadManager** class is responsible for business logic, and it acts as a mediator between **StreamlitManager** and **Graph** class. The **Graph** class contains our algorithm code and low level logic related to handling road soft deletions, restorations, computing shortest path and so on.

Libraries Used:

The following Python libraries are used in this project to accomplish its tasks. The **Shapely** library is used for click detection on the map. Specifically, the **MultiLineString** class represents a road's geometric polyline as a Shapely geometry object, and the `distance` method of the **Point** class computes the shortest distance between a user's click point and each visible road's polyline in `StreamlitManager.get_road_id_by_selection` in `streamlit_manager.py`. Since the `StreamlitManager.get_road_id_by_selection` is only called when some polyline is clicked, the job now is to figure out which line was clicked, which we do by computing distance to each polyline through Shapely and returning the road associated with the polyline with least distance to click. The **orjson** library's `orjson.loads` function in `fast_geo_json_file_loader.py` parses the big road network geojson file efficiently into a Python dictionary (a lot faster than json library does), which is then used by `RoadManager.fetch_data_and_build_graph` method in `road_manager.py` to build the graph.

The **Streamlit-Folium** library's `st_folium` function renders the interactive map and captures user clicks, enabling bidirectional communication between the map and the app. Finally, the **Folium** library's **Map** class creates the base map of Ontario, and its **GeoJson** class draws each road segment using the `geometry` attribute of each **Road** object.

For UI, streamlit (on top of which streamlit.folium is built) library's `text` method is used to show the length of new path, length of previous shortest path before road removal besides the map (Snowflake Inc., "Basic Concepts of Streamlit"; "St.text").

An optimized and modified version of Dijkstra's algorithm is implemented in `Graph.compute_shortest_path` in `graph.py`. Unlike the standard algorithm which returns only one shortest path, this version collects all shortest paths of equal length using **PathTree** class (ignoring decimals for reasons discussed earlier). The result is returned as a **ShortestPathResult** dataclass containing the shortest length and all equal shortest paths. If no path exists between the two vertices, **None** is returned to indicate disconnection between source and target vertices.

Apart from the shortest path algorithm, the program performs additional computations on the data. To avoid performance issues from rendering over one million road segments at once, `StreamlitManager._update_visible_roads_by_bounds` in `streamlit_manager.py` filters roads based on the two criteria described in bullet point 4 above. The average road length, stored as `_road_average_length`, is computed during initialization and used as part of this threshold calculation.

Instructions for Running the Program

To run this program, follow these steps:


1. Download the `ontario_road_network.geojson` file shared via one drive shared link- U of T One Drive Link Click Here (requires University of Toronto login), unzip it, and place `ontario_road_network.geojson` at the project's root directory. Plan B: In case if the link does not work, check the email `csc111-2026-01@cs.toronto.edu` for an email from `elias.alhazzaa@mail.utoronto.ca`
2. Ensure Python 3.13.5 is installed and install all required libraries using: `pip install -r requirements.txt` .
3. Open your terminal and run the following command from the root of the project: `streamlit run main.py` .

Important Note: If you are using Linux, don't run this project from **root directory of Linux**, as Streamlit does not work from there (Snowflake Inc., "Basic Concepts of Streamlit"). Note: Do not attempt to run the program by clicking the "run" button in your IDE. Streamlit applications must be launched via the terminal command mentioned above to function correctly.

The next section describes what can be expected to be seen:


2 Results

The initial screen which opens up after starting the app will load data and map, and then it would show the map at initial settings as shown below. This is the University of Toronto area along with a lot of other surrounding places,



Now go to St. George Street near Willcocks and Ursula Franklin, and click on St. George Street going from south to north (the positive road, not the one to the left of this road going the opposite way) as shown below. After clicking it, it would turn red as shown in the image indicating that it has been selected to be removed. Don't click on the one close to it: even though everything would work fine, but you would lose the track of this demo. You can also click reset button anytime you click something by mistake that you did not intend to click.

Note: Throughout this demo, you might feel as if some polylines have disappeared. This happens because only polylines in your previous bounds and zoom are visible, and the new polylines don't get added despite your current zoom and map position unless you press the Reload Polylines button to reload them as per your new bounds and zoom level. This is intentional for performance reasons, otherwise frequent updates to map and slowness of Python would degrade performance a lot.



Press Compute Shortest Path button, and the shortest path between St. George at Ursula Franklin junction and St. George at Willcocks Street junction would get computed with the previously selected portion being considered as deleted (this is simulating that road segment having a road closure). The shortest path computed would get displayed on the map in green colour, while the length of previous shortest path before closure and current shortest path after closure would be shown on the right. If you hover your mouse over the green road segments and add all the distances up, you will get the number on the right (but careful not to miss adding a road segment, a different road segment begins where the road meets another road, so in the image below, there are 4 road segments not 3, with the one on the left having two parts due to road meeting in the middle).

Deploy
⋮


Positive roads go south to north or west to east

Negative roads go north to south or east to west

Blue roads are positive roads and black roads are negative roads

Red roads are removed roads

Green roads are the ones that represent the shortest path when asked.



Roads Currently Selected: 1

Reload Polyline

Compute Shortest Path

Reset

Shortest paths computed successfully.

There is in total 1 shortest path.

Displaying path number 1 out of 1 path

Previous length of shortest path before removal of road segment was 189.0 meters.

Now, after removal, the length of shortest path is 478.0 meters.

Please note that if you zoom in too much and then perform some action, you will have to zoom out and press reload polyline to see polyline after your action is done. For instance, if you zoom in to remove a polyline, and then

Now, we will see a demonstration of this with multiple path segments (major addition apart from multiple shortest paths). Click on St. George Street North side (positive road) again, but this time on the segment which goes from Willcocks to Harbord Street. It will get removed (meaning it will turn red in colour). After you it removed as per its colour, press Compute Shortest Path button, and zoom out if necessary. The green road segments all together form the shortest path from St. George at Ursula Franklin to St. George at Harbord Street after the road segments from from Ursula Franklin to Harbord at St, George street have been removed. The result should be similar to the image below depending on your zoom level and map position.

Deploy
⋮

Positive roads go south to north or west to east

Negative roads go north to south or east to west

Blue roads are positive roads and black roads are negative roads

Red roads are removed roads

Green roads are the ones that represent the shortest path when asked.

Roads Currently Selected: 2

Reload Polylines

Compute Shortest Path

Reset

Shortest paths computed successfully.

There is in total 1 shortest path.


Displaying path number 1 out of 1 path

Previous length of shortest path before removal of road segment was 423.0 meters.



Now, after removal, the length of shortest path is 687.0 meters.

Please note that if you zoom in too much and then perform some action, you will have to zoom out and press reload polylines after your action is done. For instance, if you zoom in to

Press reset button to take every road to original state again. Now, we will see a demo for multiple shortest paths (our first major project enhancement). Scroll to Spadina at Dupont and press Reload Polylines if necessary. The image below shows that location.




Click on Spadina Road that goes to Dupont Street (positive road) to remove it. Now click compute shortest path button. This will now show that there are two different equivalent shortest paths between the two junctions, but only one of them will get displayed to make it easy to follow. Click on Show Another Shortest Path, and the second shortest path will be displayed. Pressing the button furthermore will alternate between these two paths shown in the two images below (as there are only two equivalent shortest paths).




Press reset button. Now, we will explore our second last demo: the demo of invalid selection. Click on Spad-


ina Road, and after it is removed, then click on Walmer Road which is parallel to it and does not join it directly. Then, click on Compute Shortest Path, but this time a shortest path will not come. This is because our `Graph.check_is_valid_road_selection` method will return `None`, as these we have two free start points and two free end points now violating our definition of validity. This means that we cannot decide the two vertices between which to run our `Graph.compute_shortest_path` method, and hence `compute_shortest_path` method is not invoked this time due to road sequence validity test failure.



Our last demo will be about what our app does when removing a road segment leads to junctions getting disconnected from each other. Scroll to Queen's Park area and click Reload Polylines if necessary. On the road id shown below in the tooltip (which comes up if one hovers over a road), click to remove it.



After the segment has been removed, press Compute Shortest Path method. The result would be that there is no shortest path, as removing the road leads to the two junctions getting disconnected. This can be seen directly from the map very easily as well. The red segment is connected only to black segments which are all negative roads, meaning they go from north to south in this case. Taking away the road path and no connectivity



A summary of interactions: Clicking on button Reload Polylines: gets polylines as per your new bounds, map position and zoom level.

Selecting a road/polyline on map: it gets selected and marked for deletion if Compute Shortest Path button is pressed and the selection of roads overall is valid.

Clicking on Compute Shortest Path button: if the selection of roads is not valid, it sets state to `InvalidRoadSelectionsState` instance so as to alert the user about this. If it is valid, it computes shortest path by calling `RoadManager.get_shortest_path` method (which then calls `Graph.compute_shortest_path` method) to get the path before removal of road segments, and then calls `RoadManager.remove_road_and_get_path` to remove roads (soft deletion) and gets all shortest paths between start and end junctions decided by `RoadManager.check_removability` method from sequence of selected roads. It then sets state as `JunctionDisconnectedState` if removal of road led to the two junctions being disconnected, otherwise it sets state to instance of `ShortestPathSuccessState` with relevant details so display method of `StreamlitManager` class can render shortest path and show previous shortest path length without road removal, shortest path length after road removal.

Show Another Shortest Path Button: Only appears if there are more than one equivalent shortest paths, it calls `StreamlitManager.shift_to_next_shortest_path` method, which restores the roads currently in shortest path to normal colour and makes road in the next shortest path (next meaning the equivalent shortest path next in list, not next as in ascending) as per shortest path colouring. This helps user understand different equivalent shortest paths on the map that are equivalent without all of them being displayed together, which would have led to confusion.

Reset button: resets all roads to normal state (makes their colour as their default colour) and initialized other variables again so that the app can be used repeatedly in one session.

This demo covered all features this app has: computing and displaying shortest path after removal of one road segment, after removal of multiple road segments, computing multiple shortest paths of same length, reset feature, lazy loading of polylines for performance reasons, alerting if junctions get disconnected after road removal and flagging invalid sequence of roads. Now, the reader is free to experiment as the reader of this report likes with any of these features wherever the reader wants on the map.

Changes from Proposal

For the final submission, we have created a better way to achieve the original proposed goal of building an app that allows an official to track how a road closure would affect traffic. We had two major changes to our original proposal in terms of the ways to better achieve that goal by incorporation of feedback.

The biggest change done was to use Dijkstras to give all shortest paths, not just one shortest path. This required a study of the proof of Dijkstras algorithm, coming up with own proof of a modified Dijkstras algorithm we wrote where we had to show that when a vertex is added in `all_shortest_paths_computed`, all possible shortest paths to it have been computed already, and implementing it in $O(V+E \log V)$ upper bound time. Its performance was improved by using Python's `heapq` module's `heap` (which are also a type of tree) as priority queue, using list/tuples instead of dataclasses after empirically observing slowness of dataclasses, and finding a way to capture all shortest paths without passing around a copy of list which would also have slowed this method down a lot. Instead, a reference to previous vertex from which shortest distance came was kept in distance dictionary (in `Graph.compute_shortest_path` method) along with the road id it came from (as one vertex might have multiple roads to current vertex), and then using `PathTree` (a tree) to model path back from target to source where target was the root and were the leaves (discussed the why in more detail before), and then constructing a list of list using this tree to give back a list of shortest paths. As discussed in introduction section, showing multiple paths allows the person in charge to get a relatively better view of how people will use roads. This algorithm also ignores decimals due to what was mentioned in introduction section.

Our second major change from the proposal of our project was that instead of removing one road segment and finding any shortest path from the start point to the endpoint, we now give the user the choice to remove multiple segments. If the user chooses to compute shortest path, then in `Graph.check_is_valid_road_selection` method, we check if sequence of roads is valid or not with a custom mini algorithm. We check validity by checking that there is at least one free start point (free meaning the start point is not the end point of another road) and at least one free end point (free meaning in this case that the end point is not the start point of another Road), and if there are more than one of roads with free start points and/or free end points, then we check if all such roads have the same start/end point respectively. This algorithm, if it decides the selection of roads as 'valid,' gives us a clean start junction id and end junction id between which we can compute the shortest distance. This algorithm allows the users to select multiple road segments, but in a manner that the selected road segments are not arbitrarily floating everywhere, hence making the idea of shortest path computation between two points meaningless. Since there was the possibility of by mistake selecting 'invalid' paths, we also created a reset button. Before this, Dijkstras algorithm would have run after the road segment was actually removed from the graph, but now, it would have to be soft deleted and not hard deleted to allow for a reset (a reset would include all deleted roads being restored, see `StreamlistManager.reset` method and `RoadManager.restore_removed_roads` and `Graph.restore_removed_roads` that deal exactly with this). Hence, this change also forced us indirectly to add soft deletion for road (which is a directed weighted edge), and then to modify Dijkstras in a way that it ignores the soft deleted ones. This was not possible without actually understanding how the algorithm works, and hence this compounded the complexity of soft deletion.

3 Discussion

The goal to create an app that allows a transportation person in charge to analyze how road closure affects traffic has been achieved to a degree. The following two example make it clear that a person in charge using this app can get an idea of effects of road closure on motorists in Ontario, just with a caveat discussed below. For road segments like the one near Queen's Park in our demo section, removing/closing which leads to the two junctions that the road segment used to connected getting disconnected, this program helps the person in charge immediately comprehend the vitality of that route, and hence that individual can then take steps to mitigate the closure to lessen its deep effects on the traffic where they cannot reach their destinations if they start at one of those junctions and have to go to the other junction. As a second example, consider the two equivalent shortest paths given on Spadina Dupont example. Both of them had length 504 meters, while before closure, the length was 309 meters. This is a trivial distance to cover for a motorist, but looking at both paths at the map, both of them have 2 turns, and hence despite less increase in difference, there would be considerable consequences for the motorists in this area due to swelling of intersections while cars turn, which leads to traffic slowdown. In more general terms, the combination of length of shortest map and displaying different shortest paths on map helps understand both the increase in distance and increase in number of turns for motorists, and hence gives a comprehensive picture for road closures.

However, despite this helping to achieve the goal to a degree, it does not give the complete picture. The complete picture would have been given by using Google Map's APIs to calculate average congestion and roads, and then modelling how the congestion now would increase over the roads above average congestion as a road segment is

closed. This is because if the closure of shortest path leads users to use a road that was not used a lot previously anyways, this would mean that traffic flow would not have as much friction as anticipated by the map.

Another way to improve is to use Google Map's APIs to use average time at specific time of day instead of road distance to determine which road users would take. This is from a motorist centric point of view. Motorists who use map applications don't take roads that are the shortest in distance, and rather take roads which are the shortest in time, and this is something in which the API would have helped in achieving. Thus, a Dijkstras over average time a road segment takes from junction A to junction B would have given shortest path that the user is likely to take, and this means just some slight tweaks to this project can make it more comprehensive (assuming one is willing to pay the dollars for using Google Maps API), as we are replacing just length with time now but other logic to a big degree remains the same.

It can therefore be said that users who don't rely on apps like Google Maps that prioritize time over distance, and rather rely over the length of the road they take, then this analysis covers their behaviour well.

Other than that, another limitation of this project was that roads that ran in both directions, they had their polylines at the same coordinates. This meant that a slight translation had to be done to the polyline depending on if it was of a negative road or positive road to separate bidirectional road's opposite lanes from each other. However, this slight nudge was not uniform everywhere, and sometimes it led to too much of translation whereas other times it led to too little translation of the polyline (despite translation being done using a constant), leading to confusion over which roads are located where in some cases. However, this did not affect the overall goal of this project, but just affected efficiency of the one who uses this, or maybe even slight mistakes if the user of this app is not experienced in using this app that much. This cannot be improved, as Ontario data source officially does not differentiate between polylines of the opposite lands of a bidirectional road segment.

References

Works Cited:

GeeksforGeeks. "Dijkstra's Algorithm to Find Shortest Paths from a Source to All." GeeksforGeeks, 25

Nov. 2012, www.geeksforgeeks.org/dsa/dijkstras-shortest-path-algorithm-greedy-algo-7/.

Gillies, Sean. "Shapely.MultiLineString — Shapely 2.1.2 Documentation." Readthedocs.io, 2025, shapely.readthedocs.io/en/stable/reference/shapely.MultiLineString.html.

Legislative Assembly of Ontario. "More about Ontario — Legislative Assembly of Ontario." Ola.org,

2024, www.ola.org/en/visit-learn/teach-learn-play/about-ontario/more-about-ontario.

Ontario Ministry of Natural Resources. "Ontario Road Network (ORN) Road Net Element." *Gov.on.ca*,

2020, geohub.lio.gov.on.ca/datasets/mnrf::ontario-road-network-orn-road-net-element/about.

—. ORN Road Net Element - User Guide-EN. 2025, Ontario Road Network - Road Net Element - User Guide (Word).

Snowflake Inc. "Basic Concepts of Streamlit - Streamlit Docs." Streamlit.io, 2026, docs.streamlit.io/get-started/fundamentals/main-concepts.

—. "St.text - Streamlit Docs." Streamlit.io, 2024, docs.streamlit.io/develop/api-reference/

`text/st.text.`

Story, Rob. “PolyLine — Folium 0.1.Dev1+G3ccc47c Documentation.” Python-Visualization.github.io, `python-visualization.github.io/folium/latest/user_guide/vector_layers/polyline.html`.

—. “Getting Started — Folium 0.1.Dev1+G9ab63f7 Documentation.” Python-Visualization.github.io, `python-visualization.github.io/folium/latest/getting_started.html`.

The Weather Network. “BURIED: Southern Ontario Digs out of High-Impact Snowstorm.” The Weather Network, 15 Jan. 2026, `www.theweathernetwork.com/en/news/weather/severe/buried-major-snowstorm-brings-road-closures-flight-cancellations-dvp-closed`. Accessed 4 Mar. 2026.

Zwitch, Randy. “Callbacks.” Streamlit.app, 7 July 2023, `folium.streamlit.app/callbacks`.